

assignment

April 1, 2024

```
[1]: import pandas as pd
```

```
[2]: df = df = pd.read_csv('../Data/automobileEDA.csv')
df.head()
```

```
[2]:
```

	symboling	normalized-losses	make	aspiration	num-of-doors	\
0	3	122	alfa-romero	std	two	
1	3	122	alfa-romero	std	two	
2	1	122	alfa-romero	std	two	
3	2	164	audi	std	four	
4	2	164	audi	std	four	

	body-style	drive-wheels	engine-location	wheel-base	length	...	\
0	convertible	rwd	front	88.6	0.811148	...	
1	convertible	rwd	front	88.6	0.811148	...	
2	hatchback	rwd	front	94.5	0.822681	...	
3	sedan	fwd	front	99.8	0.848630	...	
4	sedan	4wd	front	99.4	0.848630	...	

	compression-ratio	horsepower	peak-rpm	city-mpg	highway-mpg	price	\
0	9.0	111.0	5000.0	21	27	13495.0	
1	9.0	111.0	5000.0	21	27	16500.0	
2	9.0	154.0	5000.0	19	26	16500.0	
3	10.0	102.0	5500.0	24	30	13950.0	
4	8.0	115.0	5500.0	18	22	17450.0	

	city-L/100km	horsepower-binned	diesel	gas
0	11.190476	Medium	0	1
1	11.190476	Medium	0	1
2	12.368421	Medium	0	1
3	9.791667	Medium	0	1
4	13.055556	Medium	0	1

[5 rows x 29 columns]

```
[5]: x = df.drop("price", axis=1)
y = df['price']
```

```
[6]: from sklearn.model_selection import train_test_split

[48]: X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.2,
↳ random_state=42)

[49]: to_select = ['horsepower', 'curb-weight', 'engine-size', 'highway-mpg']
x_train = X_train[to_select].values
x_test = X_test[to_select].values

[3]: from sklearn.linear_model import LinearRegression
model = LinearRegression()

[52]: model.fit(x_train, y_train)

[52]: LinearRegression()

[58]: y_pred_train = model.predict(x_train)
y_pred_test = model.predict(x_test)

[59]: model.intercept_

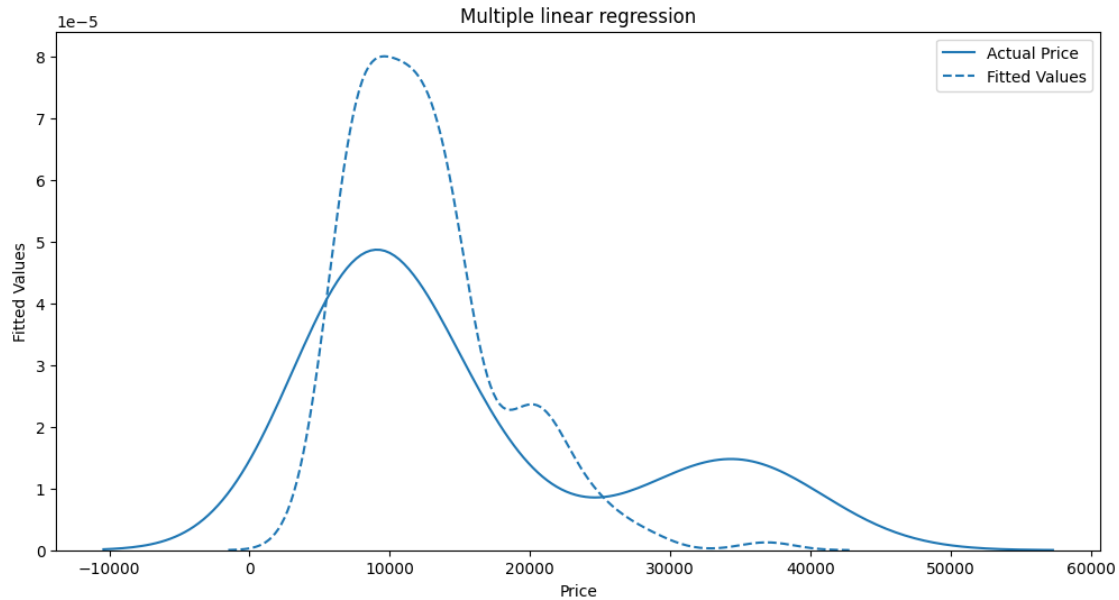
[59]: -8009.241208155732

[61]: model.coef_

[61]: array([ 25.9407567 ,   3.7778736 ,  83.14757024, -65.72539722])

[4]: import matplotlib.pyplot as plt
import seaborn as sns

[41]: plt.figure(figsize=(12, 6))
ax1 = sns.kdeplot(y_test, color='r', label='Actual Price')
sns.kdeplot(y_pred_train, color='b', label='Fitted Values', ax=ax1,
↳ linestyle='--') #Color is not changed so linestyle is changed
plt.title("Multiple linear regression")
plt.xlabel("Price")
plt.ylabel("Fitted Values")
plt.legend()
plt.show()
plt.close()
```



0.1 Mean Square Error (MSE)

The Mean Square Error (MSE) is a common loss function used in regression problems to measure the average of the squares of the errors—that is, the average squared difference between the estimated values and the actual value. The MSE is always non-negative, and values closer to zero are better. The formula for calculating MSE is given by:

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

Where: - n is the number of observations, - y_i is the actual value of the i -th observation, - \hat{y}_i is the predicted value for the i -th observation.

```
[7]: x= df['horsepower'].values.reshape(-1, 1)
y = df['price'].values.reshape(-1, 1)

X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.2,
↳random_state=42)
model = LinearRegression()
model.fit(X_train, y_train)
```

```
[7]: LinearRegression()
```

```
[8]: y_pred_train = model.predict(X_train)
y_pred_test = model.predict(X_test)
```

```
[10]: from sklearn.metrics import mean_squared_error
```

```
[11]: mse_train = mean_squared_error(y_train, y_pred_train)
mse_test = mean_squared_error(y_test, y_pred_test)

print(f"Train Mean Squared Error: {mse_train}")
print(f"Test Mean Squared Error: {mse_test}")

print('\n')
print("Coefficients:", model.coef_)
print("Intercept:", model.intercept_)
```

Train Mean Squared Error: 16320010.468020136

Test Mean Squared Error: 46093655.049259596

Coefficients: [[151.88933994]]

Intercept: [-2907.17676069]

0.2 MSE using numpy

```
[12]: import numpy as np
```

```
[13]: mse=np.mean((y_train - y_pred_train) ** 2)
mse
```

```
[13]: 16320010.468020136
```

0.3 MSE using sklearn


```
[66]: from sklearn.metrics import mean_squared_error
```

```
[68]: mse_train = mean_squared_error(y_train, y_pred_train)
mse_test = mean_squared_error(y_test, y_pred_test)
mse_train, mse_test
```

```
[68]: (8719832.96736771, 28916467.462179784)
```

0.4 Visualization

```
[40]: plt.figure(figsize=(12,6))
ax1 = sns.kdeplot(y_test, label='Actual Price')
sns.kdeplot(y_pred_train, label= 'Predicted Price', ax=ax1,
linestyle='--')#Color is not changed so linestyle is changed
plt.xlabel('Actual Price')
plt.ylabel('Predicted Price')
plt.title('Actual vs Predicted Price KDE Plot')
plt.legend()
plt.show()
```



0.5 Mean Absolute Error (MAE)

The Mean Absolute Error (MAE) is a metric used to measure the average magnitude of the errors between pairs of predictions and actual outcomes, without considering their direction. It's the mean over the test sample of the absolute differences between predicted and actual values. MAE provides a straightforward indication of average error magnitude in the same units as the data. The MAE formula is:

$$\text{MAE} = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$

Where: - n is the number of observations, - y_i is the actual value for the i -th observation, - \hat{y}_i is the predicted value for the i -th observation.

MAE is particularly useful for understanding the average error magnitude directly in the output variable's units.

```
[15]: from sklearn.metrics import mean_absolute_error
```

```
[16]: mse_train = mean_absolute_error(y_train, y_pred_train)
mse_test = mean_absolute_error(y_test, y_pred_test)
mse_train, mse_test
```

```
[16]: (2877.7586831038134, 4809.393573026446)
```

0.6 Root Mean Square Error (RMSE)

The Root Mean Square Error (RMSE) is another metric used to measure the average magnitude of the error. It's the square root of the average of squared differences between the predicted values and actual values. Compared to MAE, RMSE gives a relatively high weight to large errors. This means RMSE is more sensitive to outliers than MAE. The RMSE formula is:

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

Where: - n is the number of observations, - y_i is the actual value of the i -th observation, - \hat{y}_i is the predicted value for the i -th observation.

RMSE is valuable when large errors are particularly undesirable and should be penalized more than smaller errors. It is also more appropriate than MAE when the error distribution is expected to be Gaussian.

```
[17]: from sklearn.metrics import root_mean_squared_error
```

```
[18]: mse_train = root_mean_squared_error(y_train, y_pred_train)
mse_test = root_mean_squared_error(y_test, y_pred_test)
mse_train, mse_test
```

```
[18]: (4039.803270955176, 6789.230814257208)
```